

Transformers for image recognition

Ahmad HARKOUS

An Vinh Lala

Valentin Dauzere-Peres

<https://arxiv.org/pdf/2010.11929>

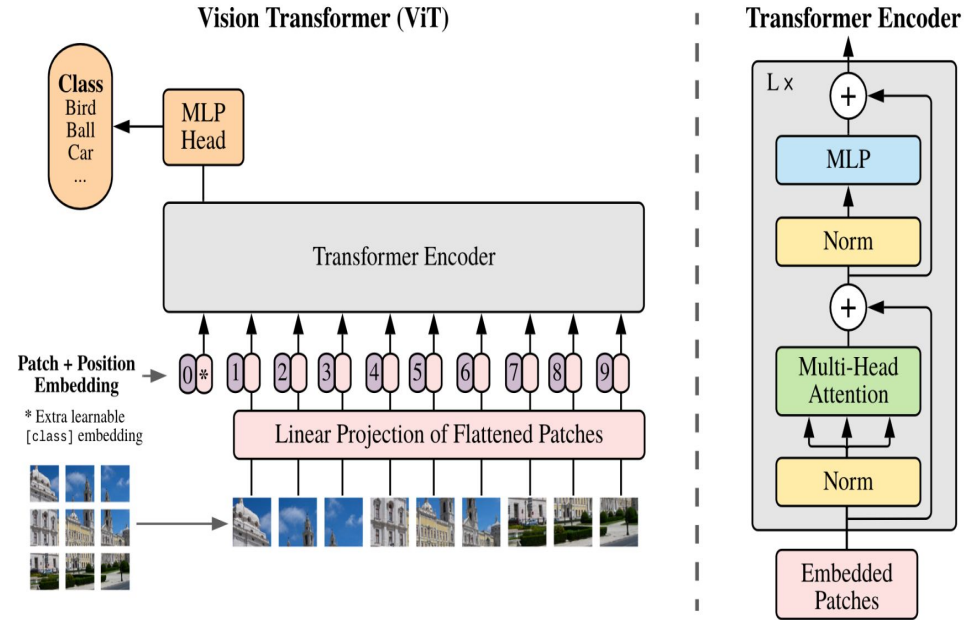
Dec 2024

Data Set

- **CIFAR-10 dataset:** 60,000
32x32 color images across 10
classes (e.g., airplanes,
automobiles, birds)
- **Training Data:** A subset of 20%
of the training set was used
- **Test Data:** The full
10,000-image test set was
retained

Model Architecture Overview

1. Patch Embeddings
2. Transformer Encoder
3. Classification Head



1 > Patch Embeddings: 3D -> 1D

> Dividing the Image into Patches

Given an input image I with dimensions $H \times W \times C$, where:

- H : Height of the image.
- W : Width of the image.
- C : Number of channels (e.g., 3 for RGB).

The image is divided into non-overlapping patches of size $P \times P$. The total number of patches N is calculated as:


$$N = \frac{H \times W}{P^2}$$

> Flattening Each Patch

Each patch of size $P \times P \times C$ is flattened into a 1D vector of size:

$$\text{Patch Size} = P^2 \cdot C$$

> Original 3D vector:



1 > Patch Embeddings: 1D -> D-dimensional

> Linear Projection

Each flattened patch is linearly projected into a D -dimensional embedding space:

$$\mathbf{e}_p = \mathbf{W} \cdot \mathbf{x}_p + \mathbf{b}$$


where:

- \mathbf{x}_p : Flattened patch (of size $P^2 \cdot C$).
- \mathbf{W} : Learnable weight matrix of size $D \times (P^2 \cdot C)$.
- \mathbf{b} : Bias term.

The result is an embedding \mathbf{e}_p of size D for each patch.

The embeddings of all N patches are concatenated to form a sequence:

$$\mathbf{E} = [\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_N]$$



1 > Patch Embeddings: D-dimensional + CLS + Position

> Prepending the [CLS] Token


$$\mathbf{E}_{\text{input}} = [[\text{CLS}], \mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_N]$$

where $\mathbf{E}_{\text{input}}$ is of size $(N + 1) \times D$.

> Adding Positional Encoding

$$\mathbf{E}_{\text{input}} = \mathbf{E}_{\text{input}} + \mathbf{P}$$

where \mathbf{P} is the positional encoding matrix of size $(N + 1) \times D$.


$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q \cdot K^T}{\sqrt{D_k}}\right) \cdot V$$
$$\text{FFN}(x) = \text{GELU}(x \cdot \mathbf{W}_1 + \mathbf{b}_1) \cdot \mathbf{W}_2 + \mathbf{b}_2$$

Layer normalization (LN) is applied before every block, and residual connections after every block of both the MHSA and FFN(MLP).

3 > Classification Head

$$\mathbf{z}_{\text{CLS}} = \text{Transformer Output}[\text{CLS}]$$

$$\text{Logits} = \mathbf{z}_{\text{CLS}} \cdot \mathbf{W}_{\text{head}} + \mathbf{b}_{\text{head}}$$

MLP HEAD

final output is a vector of size num_classes, containing the predicted class logits.


$$\mathbf{x} = [x_1, x_2, x_3, x_4, x_5, x_6, x_7, x_8, x_9, x_{10}]$$

Simplifications


- Using fewer Transformer layers (6 layers instead of the original 12)
- Reducing embedding dimensions D to 64
- Limiting self-attention heads to 4
- Fixed learning rate and fewer epochs for tuning,

Performances



Training Loss



Test Accuracy



Article Performances



Questions ?